

Project 1: Canny's Edge Detector

Cat Mai

Requirements

1. Being able to compile Python code
2. Install Pillow and numpy, on some command line
 - a. pip install Pillow
 - b. pip install numpy
3. Change img_path in `main.py`

```
if __name__ == "__main__":  
  
    img_path = "Barbara.bmp"  
    smooth = gaussian_smooth(img_path)
```

Source Code

```
1  import PIL  
2  from PIL import Image  
3  import numpy as np  
4  from matplotlib import pyplot as plt  
5  
6  
7  def convolve(img, mask):  
8      # img and mask are both np arrays  
9  
10     mask_width, mask_height = mask.shape  
11     result_width, result_height = img.shape  
12     result = np.zeros(img.shape)  
13  
14     for row in range(result_width-mask_width):  
15         for col in range(result_height-mask_height):  
16  
17             img_window = img[row:row+mask_width, col:col+mask_height]  
18             result[row,col] = np.sum(np.multiply(img_window, mask))  
19  
20     return result  
21  
22
```

```

22
23 ▼ def gaussian_smooth(img_path):
24
25     # perform gaussian smoothing
26 ▼     gaussian_mask = [
27         [1,1,2,2,2,1,1],
28         [1,2,2,4,2,2,1],
29         [2,2,4,8,4,2,2],
30         [2,4,8,16,8,4,2],
31         [2,2,4,8,4,2,2],
32         [1,2,2,4,2,2,1],
33         [1,1,2,2,2,1,1]
34     ]
35
36     mask = np.array(gaussian_mask)
37     img = np.asarray(Image.open(img_path))
38     result = convolve(img, mask)
39
40     # normalize
41     return result/np.sum(gaussian_mask)
42
43

```

```

44 def get_gradients(img):
45
46     # Four masks provided
47     g0 = np.array([
48         [-1, 0, 1],
49         [-2, 0, 2],
50         [-1, 0, 1]
51     ])
52
53     g1 = np.array([
54         [0, 1, 2],
55         [-1, 0, 1],
56         [-2, -1, 0]
57     ])
58
59     g2 = np.array([
60         [1, 2, 1],
61         [0, 0, 0],
62         [-1, -2, -1]
63     ])
64
65     g3 = np.array([
66         [2, 1, 0],
67         [1, 0, -1],
68         [0, -1, -2]
69     ])
70
71     # get abs of responses from the four masks
72     m0 = np.abs(convolve(img, g0))
73     m1 = np.abs(convolve(img, g1))
74     m2 = np.abs(convolve(img, g2))
75     m3 = np.abs(convolve(img, g3))
76
77     # magnitude is the maximum of the four responses, divided by 4
78     gradient_magnitude = np.maximum.reduce([m0, m1, m2, m3])/4
79
80
81     # quantized angle equals to the index of the mask that produces the maximum response
82     quantized_angles = np.zeros(gradient_magnitude.shape)
83     quantized_angles_width, _ = quantized_angles.shape
84
85     for row in range(quantized_angles_width):
86         row_array = np.array([m0[row], m1[row], m2[row], m3[row]])
87         quantized_angles[row] = np.argmax(row_array, axis=0)
88
89     return gradient_magnitude, quantized_angles
90

```

```

92  def nms(magnitude, angles):
93      # perform non-maxima suppression
94
95      # indices of the two neighbors to do the comparison
96      neighbors = {
97          0: [(0,-1), (0,1)],
98          1: [(1,-1), (-1,1)],
99          2: [(-1,0), (1,0)],
100         3: [(-1,-1), (1,1)]
101     }
102
103     nms_magnitude = np.zeros(magnitude.shape)
104     width, height = magnitude.shape
105
106     # get neighbors for each pixels and do nms comparison
107     for row in range(width):
108         for col in range(height):
109
110             try:
111                 # neighbor 1
112                 n1_row, n1_col = neighbors[angles[row][col]][0]
113                 n1_mag = magnitude[row+n1_row][col+n1_col]
114
115                 # neighbor 2
116                 n2_row, n2_col = neighbors[angles[row][col]][1]
117                 n2_mag = magnitude[row+n2_row][col+n2_col]
118
119                 curr_mag = magnitude[row][col]
120                 if curr_mag < n1_mag or curr_mag < n2_mag:
121                     nms_magnitude[row][col] = 0
122                 else:
123                     nms_magnitude[row][col] = curr_mag
124
125             except IndexError:
126                 # the pixel doesn't have two neighbors, do nothing
127                 pass
128
129     return nms_magnitude
130


```

```


131
132 ▼ def threshold(nms_magnitude):
133
134     # perform simple thresholding for 25, 50, 75th percentile
135
136     vals = [v for v in nms_magnitude.flatten() if v]
137     t25, t50, t75 = np.percentile(vals, [25, 50, 75])
138
139     edgemap25 = np.zeros(nms_magnitude.shape)
140     edgemap50 = np.zeros(nms_magnitude.shape)
141     edgemap75 = np.zeros(nms_magnitude.shape)
142     width, height = nms_magnitude.shape
143
144     for row in range(width):
145     ▼     for col in range(height):
146
147         mag = nms_magnitude[row][col]
148     ▼         if mag >= t25:
149             edgemap25[row][col] = 255
150     ▼         if mag >= t50:
151             edgemap50[row][col] = 255
152     ▼         if mag >= t75:
153             edgemap75[row][col] = 255
154
155     return edgemap25, edgemap50, edgemap75
156
157 ▼ def histogram(nms_magnitude):
158     vals = [v for v in nms_magnitude.flatten() if v]
159     _ = plt.hist(vals, bins='auto')
160     plt.show()
161
162
163 ▼ if __name__ == "__main__":
164
165     img_path = "Peppers.bmp"
166     smooth = gaussian_smooth(img_path)
167     mag, angles = get_gradients(smooth)
168     nms = nms(mag, angles)
169     edgemap25, edgemap50, edgemap75 = threshold(nms)
170     histogram(nms)
171

```


Histograms



Goldhill_histogram.png



Peppers_histogram.png



Barbara.bmp

After Gaussian smoothing



Normalized mag image



After nms



25th Percentile - 2.89



50th Percentile - 6.64



75th Percentile - 14.7



Goldhill.bmp

After Gaussian smoothing



Normalized mag image



After nms



25th Percentile - 3.56



50th Percentile - 7.2



75th Percentile - 14.4



Peppers.bmp

After Gaussian smoothing



Normalized mag image



After nms



25th Percentile - 1.69



50th Percentile - 3.59



75th Percentile - 10.56

